**ONLINE** Same question and retrieved evidence

## Prompted LLM selector

Read evidence, then choose a plan

**Input**
Question + evidence + passage count k

**LLM judgment**
Select E / C / R

**Execute plan**
Selected E / C / R

## Self-verification cascade

Generate a candidate, then check it

**Candidate answer**
Generate first

**Self-verification**
Question, evidence and candidate answer

**Keep or escalate**
Execute if escalated

## Decider (Ours)

Predict execution gains before generation

**Input**
Question + evidence + passage count k

**One 4B pass**
Predict gains with learned heads

**Calibrated choice**
Select E / C / R, then execute

**OFFLINE** Execution outcomes train the planner; validation calibrates the policy

**Outcome traces**
Execute E / C / R on common evidence

**Fit the planner**
Quality + gain targets; LoRA + 5-output head

**Calibrate policy**
Validation profiles and budget thresholds